

# Datenbank Architektur für Fortgeschrittene

## Ausarbeitung 1: Anfrageverarbeitung

Daniel Gürber  
Stefan Eggenschwiler

02.05.2013

# Inhaltsverzeichnis

<b>1</b>	<b>Vorbereitung</b>	<b>1</b>
1.1	Einrichten Datenbasis . . . . .	1
<b>2</b>	<b>Statistiken erheben</b>	<b>1</b>
<b>3</b>	<b>Ausführungsplan</b>	<b>1</b>
<b>4</b>	<b>Versuche ohne Index</b>	<b>2</b>
4.1	Projektion . . . . .	2
4.2	Selektion . . . . .	2
4.3	Join . . . . .	4
<b>5</b>	<b>Versuch mit Index</b>	<b>5</b>
5.1	Projektion . . . . .	5
5.2	Selektion . . . . .	5
5.3	Join . . . . .	7
<b>6</b>	<b>Quiz</b>	<b>9</b>
<b>7</b>	<b>Deep Left Join</b>	<b>11</b>
<b>8</b>	<b>Eigene SQL-Abfragen</b>	<b>11</b>

# 1 Vorbereitung

## 1.1 Einrichten Datenbasis

```
1 CREATE TABLE regions
2 AS SELECT *
3   FROM dbarc00.regions;
4
5 CREATE TABLE nations
6 AS SELECT *
7   FROM dbarc00.nations;
8
9 CREATE TABLE parts
10 AS SELECT *
11   FROM dbarc00.parts;
12
13 CREATE TABLE customers
14 AS SELECT *
15   FROM dbarc00.customers;
16
17 CREATE TABLE suppliers
18 AS SELECT *
19   FROM dbarc00.suppliers;
20
21 CREATE TABLE orders
22 AS SELECT *
23   FROM dbarc00.orders;
24
25 CREATE TABLE partsupps
26 AS SELECT *
27   FROM dbarc00.partsupps;
28
29 CREATE TABLE lineitems
30 AS SELECT *
31   FROM dbarc00.lineitems;
```

## 2 Statistiken erheben

```
1 BEGIN
2   DBMS_STATS.GATHER_TABLE_STATS('dbarc00','parts');
3 END;
```

	CUSTOMERS	LINEITEMS	NATIONS	ORDER	PARTS	PARTSUPPS	REGIONS	SUPPLIERS
Zellen	150000	6001215	25	1500000	200000	800000	5	10000
Bytes	23850000	750151875	2675	166500000	26400000	114400000	480	1440000
Blöcke	3494	109217	4	24284	3859	16650	4	220
Extends	43	186	1	95	46	88	1	17

## 3 Ausführungsplan

```
1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM parts;
```

```
1 SELECT plan_table_output
2 FROM TABLE(DBMS_XPLAN.DISPLAY('plan_table',null,'serial'));
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		200K	25M	1051 (1)	00:00:13
1	TABLE ACCESS FULL	PARTS	200K	25M	1051 (1)	00:00:13

Reflexion  
BLABLA

## 4 Versuche ohne Index

### 4.1 Projektion

```
1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM ORDERS;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	158M	6610 (1)	00:01:20
1	TABLE ACCESS FULL	ORDERS	1500K	158M	6610 (1)	00:01:20

Reflexion  
BLABLA

```
1 EXPLAIN PLAN FOR
2 SELECT o_clerk
3 FROM ORDERS;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	22M	6607 (1)	00:01:20
1	TABLE ACCESS FULL	ORDERS	1500K	22M	6607 (1)	00:01:20

Reflexion  
BLABLA

```
1 EXPLAIN PLAN FOR
2 SELECT DISTINCT o_clerk
3 FROM ORDERS;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1000	16000	6676 (2)	00:01:21
1	HASH UNIQUE		1000	16000	6676 (2)	00:01:21
2	TABLE ACCESS FULL	ORDERS	1500K	22M	6607 (1)	00:01:20

Reflexion  
BLABLA

### 4.2 Selektion

Exact Point Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey=44444;

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	111	6602 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	1	111	6602 (1)	00:01:20

Predicate Information (identified by operation id):

```

1 - filter("O_ORDERKEY"=44444)

```

## Reflexion

BLABLA

## Partial Point Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey=44444 OR o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1501	162K	6629 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	1501	162K	6629 (1)	00:01:20

Predicate Information (identified by operation id):

```

1 - filter("O_CLERK"='Clerk#000000286' OR "O_ORDERKEY"=44444)

```

## Reflexion

BLABLA

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey=44444 AND o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	111	6611 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	1	111	6611 (1)	00:01:20

Predicate Information (identified by operation id):

```

1 - filter("O_ORDERKEY"=44444 AND "O_CLERK"='Clerk#000000286')

```

## Reflexion

BLABLA

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey*2=44444 AND o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		15	1665	6615 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	15	1665	6615 (1)	00:01:20

Predicate Information (identified by operation id):

1 - filter("O\_ORDERKEY"\*2=44444 AND "O\_CLERK"='Clerk#000000286')

## Reflexion

BLABLA

## Range Query

```
1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey BETWEEN 111111 AND 222222
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		27780	3011K	6603 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	27780	3011K	6603 (1)	00:01:20

Predicate Information (identified by operation id):

1 - filter("O\_ORDERKEY"<=222222 AND "O\_ORDERKEY">=111111)

## Reflexion

Die Intervallgrösse spielt in diesem Beispiel keine Rolle.

## Partial Range Query

```
1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey BETWEEN 44444 AND 55555
5 AND o_clerk BETWEEN 'Clerk#000000130' AND 'Clerk#000000139'
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		6	666	6611 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	6	666	6611 (1)	00:01:20

Predicate Information (identified by operation id):

1 - filter("O\_ORDERKEY"<=55555 AND "O\_CLERK"<='Clerk#000000139' AND "O\_ORDERKEY">=44444 AND "O\_CLERK">='Clerk#000000130')

## Reflexion

BLABLA

## 4.3 Join

```
1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders, customers
4 WHERE o_custkey = c_custkey
5 AND o_orderkey < 100;
```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		25	6750	7555 (1)	00:01:31
* 1	HASH JOIN		25	6750	7555 (1)	00:01:31
* 2	TABLE ACCESS FULL	ORDERS	25	2775	6602 (1)	00:01:20
3	TABLE ACCESS FULL	CUSTOMERS	150K	22M	951 (1)	00:00:12

Predicate Information (identified by operation id):

```

1 - access("O_CUSTKEY"="C_CUSTKEY")
2 - filter("O_ORDERKEY"<100)

```

## Reflexion

Spielen Varianten in der Formulierung eine Rolle? //NO DIFFERENCE FOR INNER JOIN

## 5 Versuch mit Index

```
1 CREATE INDEX o_orderkey_ix ON orders(o_orderkey);
```

```
1 CREATE INDEX o_clerk_ix ON orders(o_clerk);
```

Indexgrösse	Index Name	Tabellengrösse in Bytes
60817408	o_orderkey_ix	166500000
96468992	o_clerik_ix	166500000

### 5.1 Projektion

```

1 EXPLAIN PLAN FOR
2 SELECT DISTINCT o_clerk
3 FROM ORDERS;

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1000	16000	1615 (5)	00:00:20
1	HASH UNIQUE		1000	16000	1615 (5)	00:00:20
2	INDEX FAST FULL SCAN	O_CLERK_IX	1500K	22M	1546 (1)	00:00:19

## Reflexion

BLABLA

### 5.2 Selektion

Exact Point Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM ORDERS
4 WHERE o_orderkey=44444

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	111	4 (0)	00:00:01
1	TABLE ACCESS BY INDEX ROWID	ORDERS	1	111	4 (0)	00:00:01
* 2	INDEX RANGE SCAN	O_ORDERKEY_IX	1		3 (0)	00:00:01

Predicate Information (identified by operation id):

```
2 - access("O_ORDERKEY"=44444)
```

```

1 EXPLAIN PLAN FOR
2 SELECT /*+ FULL(orders) */ *
3 FROM ORDERS
4 WHERE o_orderkey=44444

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	111	6602 (1)	00:01:20
* 1	TABLE ACCESS FULL	ORDERS	1	111	6602 (1)	00:01:20

Predicate Information (identified by operation id):

1 - filter("O\_ORDERKEY"=44444)

## Reflexion

### BLABLA

Partial Point Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey=44444 OR o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1501	162K	336 (0)	00:00:05
1	TABLE ACCESS BY INDEX ROWID	ORDERS	1501	162K	336 (0)	00:00:05
2	BITMAP CONVERSION TO ROWIDS					
3	BITMAP OR					
4	BITMAP CONVERSION FROM ROWIDS					
* 5	INDEX RANGE SCAN	O_CLERK_IX			8 (0)	00:00:01
6	BITMAP CONVERSION FROM ROWIDS					
* 7	INDEX RANGE SCAN	O_ORDERKEY_IX			3 (0)	00:00:01

Predicate Information (identified by operation id):

5 - access("O\_CLERK"='Clerk#000000286')

7 - access("O\_ORDERKEY"=44444)

## Reflexion

### BLABLA

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey=44444 AND o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	111	4 (0)	00:00:01
* 1	TABLE ACCESS BY INDEX ROWID	ORDERS	1	111	4 (0)	00:00:01
* 2	INDEX RANGE SCAN	O_ORDERKEY_IX	1		3 (0)	00:00:01

Predicate Information (identified by operation id):

1 - filter("O\_CLERK"='Clerk#000000286')

2 - access("O\_ORDERKEY"=44444)

## Reflexion

### BLABLA

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey*2=44444 AND o_clerk='Clerk#000000286';

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		15	1665	1464 (1)	00:00:18
* 1	TABLE ACCESS BY INDEX ROWID	ORDERS	15	1665	1464 (1)	00:00:18



```

|* 2 | INDEX RANGE SCAN | O_CLERK_IX | 1500 | | 8 (0)| 00:00:01 |
-----
Predicate Information (identified by operation id):
-----
1 - filter("O_ORDERKEY"*2=44444)
2 - access("O_CLERK"='Clerk#000000286')

```

## Reflexion

BLABLA

Range Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey BETWEEN 111111 AND 222222

```

```

-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
-----
| 0 | SELECT STATEMENT | | 27780 | 3011K | 932 (1)| 00:00:12 |
| 1 | TABLE ACCESS BY INDEX ROWID | ORDERS | 27780 | 3011K | 932 (1)| 00:00:12 |
|* 2 | INDEX RANGE SCAN | O_ORDERKEY_IX | 27780 | | 68 (0)| 00:00:01 |
-----
Predicate Information (identified by operation id):
-----
2 - access("O_ORDERKEY">=111111 AND "O_ORDERKEY"<=222222)

```

## Reflexion

Spielt der Intervall eine Rolle?

Partial Range Query

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders
4 WHERE o_orderkey BETWEEN 44444 AND 55555
5 AND o_clerk BETWEEN 'Clerk#000000130' AND 'Clerk#000000139'

```

```

-----
| Id | Operation | Name | Rows | Bytes | Cost (%CPU)| Time |
-----
| 0 | SELECT STATEMENT | | 6 | 666 | 27 (12)| 00:00:01 | |
| 1 | TABLE ACCESS BY INDEX ROWID | ORDERS | 6 | 666 | 27 (12)| 00:00:01 |
| 2 | BITMAP CONVERSION TO ROWIDS | | | | | | |
| 3 | BITMAP AND | | | | | | |
| 4 | BITMAP CONVERSION FROM ROWIDS | | | | | | |
| 5 | SORT ORDER BY | | | | | | |
|* 6 | INDEX RANGE SCAN | O_ORDERKEY_IX | 2780 | | 9 (0)| 00:00:01 |
| 7 | BITMAP CONVERSION FROM ROWIDS | | | | | | |
| 8 | SORT ORDER BY | | | | | | |
|* 9 | INDEX RANGE SCAN | O_CLERK_IX | 2780 | | 14 (0)| 00:00:01 |
-----
Predicate Information (identified by operation id):
-----
6 - access("O_ORDERKEY">=44444 AND "O_ORDERKEY"<=55555)
9 - access("O_CLERK">='Clerk#000000130' AND "O_CLERK"<='Clerk#000000139')

```

## Reflexion

BLABLA

## 5.3 Join

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders, customers
4 WHERE o_custkey = c_custkey;

```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	386M		17514 (1)	00:03:31
* 1	HASH JOIN		1500K	386M	24M	17514 (1)	00:03:31
2	TABLE ACCESS FULL	CUSTOMERS	150K	22M		951 (1)	00:00:12
3	TABLE ACCESS FULL	ORDERS	1500K	158M		6610 (1)	00:01:20

Predicate Information (identified by operation id):

1 - access("O\_CUSTKEY"="C\_CUSTKEY")

## Reflexion

### BLABLA

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders, customers
4 WHERE o_custkey = c_custkey
5 AND o_orderkey < 100;

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		25	6750	957 (1)	00:00:12
* 1	HASH JOIN		25	6750	957 (1)	00:00:12
2	TABLE ACCESS BY INDEX ROWID	ORDERS	25	2775	4 (0)	00:00:01
* 3	INDEX RANGE SCAN	O_ORDERKEY_IX	25		3 (0)	00:00:01
4	TABLE ACCESS FULL	CUSTOMERS	150K	22M	951 (1)	00:00:12

Predicate Information (identified by operation id):

1 - access("O\_CUSTKEY"="C\_CUSTKEY")  
3 - access("O\_ORDERKEY"<100)

## Reflexion

### BLABLA

```

1 CREATE INDEX c_custkey_ix ON customers(c_custkey);

```

```

1 EXPLAIN PLAN FOR
2 SELECT *
3 FROM orders, customers
4 WHERE o_custkey = c_custkey;

```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	386M		17514 (1)	00:03:31
* 1	HASH JOIN		1500K	386M	24M	17514 (1)	00:03:31
2	TABLE ACCESS FULL	CUSTOMERS	150K	22M		951 (1)	00:00:12
3	TABLE ACCESS FULL	ORDERS	1500K	158M		6610 (1)	00:01:20

Predicate Information (identified by operation id):

1 - access("O\_CUSTKEY"="C\_CUSTKEY")

## Reflexion

### BLABLA

Erzwingen eines Nested Loop Joins:

```

1 EXPLAIN PLAN FOR
2 SELECT /*+ USE_NL (o c) */ *
3 FROM orders o, customers c
4 WHERE o_custkey = c_custkey;

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	386M	3007K (1)	10:01:34
1	NESTED LOOPS					
2	NESTED LOOPS		1500K	386M	3007K (1)	10:01:34
3	TABLE ACCESS FULL	ORDERS	1500K	158M	6610 (1)	00:01:20
* 4	INDEX RANGE SCAN	C_CUSTKEY_IX	1		1 (0)	00:00:01
5	TABLE ACCESS BY INDEX ROWID	CUSTOMERS	1	159	2 (0)	00:00:01

Predicate Information (identified by operation id):

4 - access("O\_CUSTKEY"="C\_CUSTKEY")

## Reflexion

BLABLA

Erzwingen eines Nicht-Hash Joins:

```

1 EXPLAIN PLAN FOR
2 SELECT /*+ NO_USE_HASH (o c) */ *
3 FROM orders o, customers c
4 WHERE o_custkey = c_custkey;
```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1500K	386M		50568 (1)	00:10:07
1	MERGE JOIN		1500K	386M		50568 (1)	00:10:07
2	SORT JOIN		150K	22M	52M	6202 (1)	00:01:15
3	TABLE ACCESS FULL	CUSTOMERS	150K	22M		951 (1)	00:00:12
* 4	SORT JOIN		1500K	158M	390M	44366 (1)	00:08:53
5	TABLE ACCESS FULL	ORDERS	1500K	158M		6610 (1)	00:01:20

Predicate Information (identified by operation id):

4 - access("O\_CUSTKEY"="C\_CUSTKEY")  
filter("O\_CUSTKEY"="C\_CUSTKEY")

## Reflexion

BLABLA

## 6 Quiz

```

1 EXPLAIN PLAN FOR
2 SELECT count(*)
3 FROM parts, partsupps, lineitems
4 WHERE p_partkey=ps_partkey
5 AND ps_partkey=l_partkey
6 AND ps_suppkey=l_suppkey
7 AND ((ps_partkey = 5 AND p_type = 'MEDIUM ANODIZED BRASS')
8 OR (ps_partkey = 5 AND p_type = 'MEDIUM BRUSHED COPPER'));
```

Ausgangslage:

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	45	35577 (2)	00:07:07
1	SORT AGGREGATE		1	45		
* 2	HASH JOIN		4	180	35577 (2)	00:07:07
* 3	HASH JOIN		4	144	5872 (6)	00:01:11
* 4	TABLE ACCESS FULL	PARTSUPPS	4	36	4525 (1)	00:00:55
* 5	TABLE ACCESS FULL	PARTS	2667	72009	1052 (1)	00:00:13
6	TABLE ACCESS FULL	LINEITEMS	6001K	51M	29675 (1)	00:05:57

Predicate Information (identified by operation id):

2 - access("PS\_PARTKEY"="L\_PARTKEY" AND "PS\_SUPPKEY"="L\_SUPPKEY")  
3 - access("P\_PARTKEY"="PS\_PARTKEY")

```

filter("PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM ANODIZED BRASS' OR
      "PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM BRUSHED COPPER')
4 - filter("PS_PARTKEY"=5)

```

ps\_partkey ist restriktiv:

```

1 CREATE INDEX ps_partkey_ix ON partsups(ps_partkey);

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	45	31055 (2)	00:06:13
1	SORT AGGREGATE		1	45		
* 2	HASH JOIN		4	180	31055 (2)	00:06:13
* 3	HASH JOIN		4	144	1350 (23)	00:00:17
4	TABLE ACCESS BY INDEX ROWID	PARTSUPPS	4	36	4 (0)	00:00:01
* 5	INDEX RANGE SCAN	PS_PARTKEY_IX	4		3 (0)	00:00:01
* 6	TABLE ACCESS FULL	PARTS	2667	72009	1052 (1)	00:00:13
7	TABLE ACCESS FULL	LINEITEMS	6001K	51M	29675 (1)	00:05:57

Predicate Information (identified by operation id):

```

2 - access("PS_PARTKEY"="L_PARTKEY" AND "PS_SUPPKEY"="L_SUPPKEY")
3 - access("P_PARTKEY"="PS_PARTKEY")
   filter("PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM ANODIZED BRASS' OR "PS_PARTKEY"=5 AND
         "P_TYPE"='MEDIUM BRUSHED COPPER')

```

p\_type ist restriktiv:

```

1 CREATE INDEX p_type_ix ON parts(p_type);

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
0	SELECT STATEMENT		1	45	29885 (1)	00:05:59
1	SORT AGGREGATE		1	45		
* 2	HASH JOIN		4	180	29885 (1)	00:05:59
3	NESTED LOOPS					
4	NESTED LOOPS		4	144	180 (0)	00:00:03
5	TABLE ACCESS BY INDEX ROWID	PARTSUPPS	4	36	4 (0)	00:00:01
* 6	INDEX RANGE SCAN	PS_PARTKEY_IX	4		3 (0)	00:00:01
7	BITMAP CONVERSION TO ROWIDS					
8	BITMAP AND					
9	BITMAP OR					
10	BITMAP CONVERSION FROM ROWIDS					
* 11	INDEX RANGE SCAN	P_TYPE_IX			8 (0)	00:00:01
12	BITMAP CONVERSION FROM ROWIDS					
* 13	INDEX RANGE SCAN	P_TYPE_IX			8 (0)	00:00:01
14	BITMAP OR					
15	BITMAP CONVERSION FROM ROWIDS					
* 16	INDEX RANGE SCAN	P_TYPE_IX			8 (0)	00:00:01
17	BITMAP CONVERSION FROM ROWIDS					
* 18	INDEX RANGE SCAN	P_TYPE_IX			8 (0)	00:00:01
* 19	TABLE ACCESS BY INDEX ROWID	PARTS	1	27	180 (0)	00:00:03
20	TABLE ACCESS FULL	LINEITEMS	6001K	51M	29675 (1)	00:05:57

Predicate Information (identified by operation id):

```

2 - access("PS_PARTKEY"="L_PARTKEY" AND "PS_SUPPKEY"="L_SUPPKEY")
6 - access("PS_PARTKEY"=5)
11 - access("P_TYPE"='MEDIUM ANODIZED BRASS')
13 - access("P_TYPE"='MEDIUM BRUSHED COPPER')
16 - access("P_TYPE"='MEDIUM ANODIZED BRASS')
18 - access("P_TYPE"='MEDIUM BRUSHED COPPER')
19 - filter("P_PARTKEY"="PS_PARTKEY" AND ("PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM ANODIZED

```

l\_partkey ist restriktiv:

```

1 CREATE INDEX l_partkey_ix ON lineitems(l_partkey);

```

Id	Operation	Name	Rows	Bytes	Cost (%CPU)	Time
----	-----------	------	------	-------	-------------	------

0	SELECT STATEMENT		1	45	308	(0)	00:00:04
1	SORT AGGREGATE		1	45			
2	NESTED LOOPS						
3	NESTED LOOPS		4	180	308	(0)	00:00:04
4	NESTED LOOPS		4	144	180	(0)	00:00:03
5	TABLE ACCESS BY INDEX ROWID	PARTSUPPS	4	36	4	(0)	00:00:01
* 6	INDEX RANGE SCAN	PS_PARTKEY_IX	4		3	(0)	00:00:01
* 7	TABLE ACCESS BY INDEX ROWID	PARTS	1	27	180	(0)	00:00:03
8	BITMAP CONVERSION TO ROWIDS						
9	BITMAP AND						
10	BITMAP OR						
11	BITMAP CONVERSION FROM ROWIDS						
* 12	INDEX RANGE SCAN	P_TYPE_IX			8	(0)	00:00:01
13	BITMAP CONVERSION FROM ROWIDS						
* 14	INDEX RANGE SCAN	P_TYPE_IX			8	(0)	00:00:01
15	BITMAP OR						
16	BITMAP CONVERSION FROM ROWIDS						
* 17	INDEX RANGE SCAN	P_TYPE_IX			8	(0)	00:00:01
18	BITMAP CONVERSION FROM ROWIDS						
* 19	INDEX RANGE SCAN	P_TYPE_IX			8	(0)	00:00:01
* 20	INDEX RANGE SCAN	L_PARTKEY_IX	30		2	(0)	00:00:01
* 21	TABLE ACCESS BY INDEX ROWID	LINEITEMS	1	9	32	(0)	00:00:01

Predicate Information (identified by operation id):

```

6 - access("PS_PARTKEY"=5)
7 - filter("P_PARTKEY"="PS_PARTKEY" AND ("PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM ANODIZED
  BRASS' OR "PS_PARTKEY"=5 AND "P_TYPE"='MEDIUM BRUSHED COPPER'))
12 - access("P_TYPE"='MEDIUM ANODIZED BRASS')
14 - access("P_TYPE"='MEDIUM BRUSHED COPPER')
17 - access("P_TYPE"='MEDIUM ANODIZED BRASS')
19 - access("P_TYPE"='MEDIUM BRUSHED COPPER')
20 - access("PS_PARTKEY"="L_PARTKEY")
21 - filter("PS_SUPPKEY"="L_SUPPKEY")

```

Optimierte Abfrage:

```

1 EXPLAIN PLAN FOR
2 SELECT count(*)
3 FROM parts, partsupps, lineitems, orders
4 WHERE p_partkey=ps_partkey
5 AND ps_partkey=l_partkey
6 AND ps_suppkey=l_suppkey
7 AND l_orderkey=o_orderkey

```

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		1	35		54346 (1)	00:10:53
1	SORT AGGREGATE		1	35			
* 2	HASH JOIN		803K	26M	25M	54346 (1)	00:10:53
3	TABLE ACCESS FULL	ORDERS	1500K	8789K		6599 (1)	00:01:20
* 4	HASH JOIN		792K	21M	19M	44915 (1)	00:08:59
* 5	HASH JOIN		792K	10M	3328K	6540 (1)	00:01:19
6	TABLE ACCESS FULL	PARTS	200K	976K		1050 (1)	00:00:13
7	TABLE ACCESS FULL	PARTSUPPS	800K	7031K		4523 (1)	00:00:55
8	TABLE ACCESS FULL	LINEITEMS	6001K	85M		29675 (1)	00:05:57

Predicate Information (identified by operation id):

```

2 - access("L_ORDERKEY"="O_ORDERKEY")
4 - access("PS_PARTKEY"="L_PARTKEY" AND "PS_SUPPKEY"="L_SUPPKEY")
5 - access("P_PARTKEY"="PS_PARTKEY")

```

Reflexion  
BLABLA

## 7 Deep Left Join

## 8 Eigene SQL-Abfragen